

# Assignment2

June 8, 2017

## 1 Load the dataset



```
In [1]: import pandas as pd
        df=pd.read_csv("datasets/train.tsv",delimiter="\t")
        df_good=df[df["Label"]==1]
        df_bad=df[df["Label"]==2]
        print("Done reading")
```


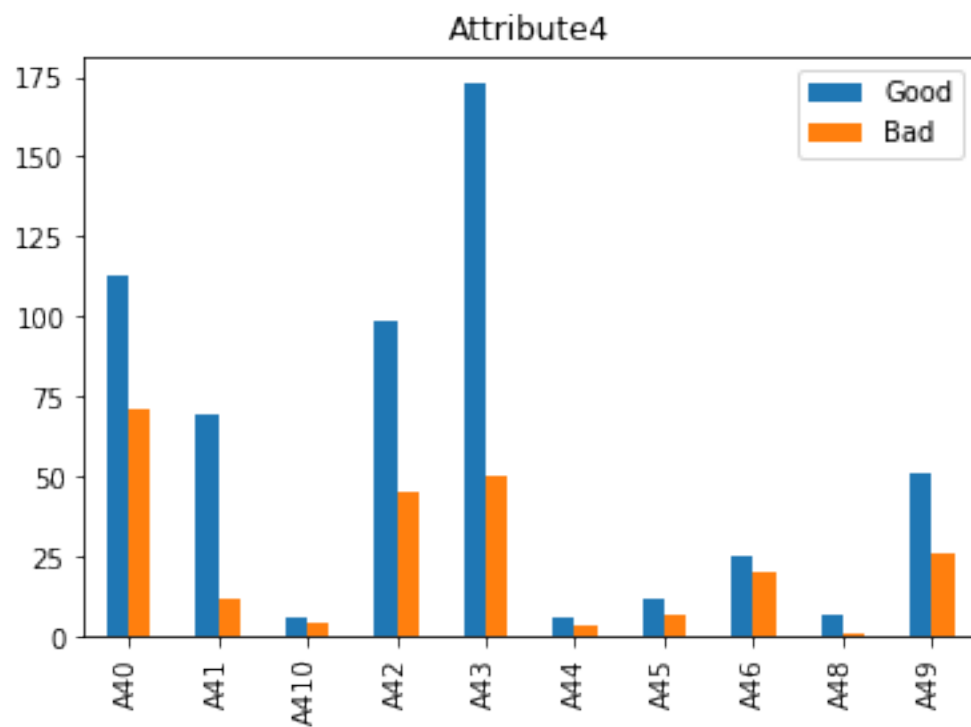
Done reading

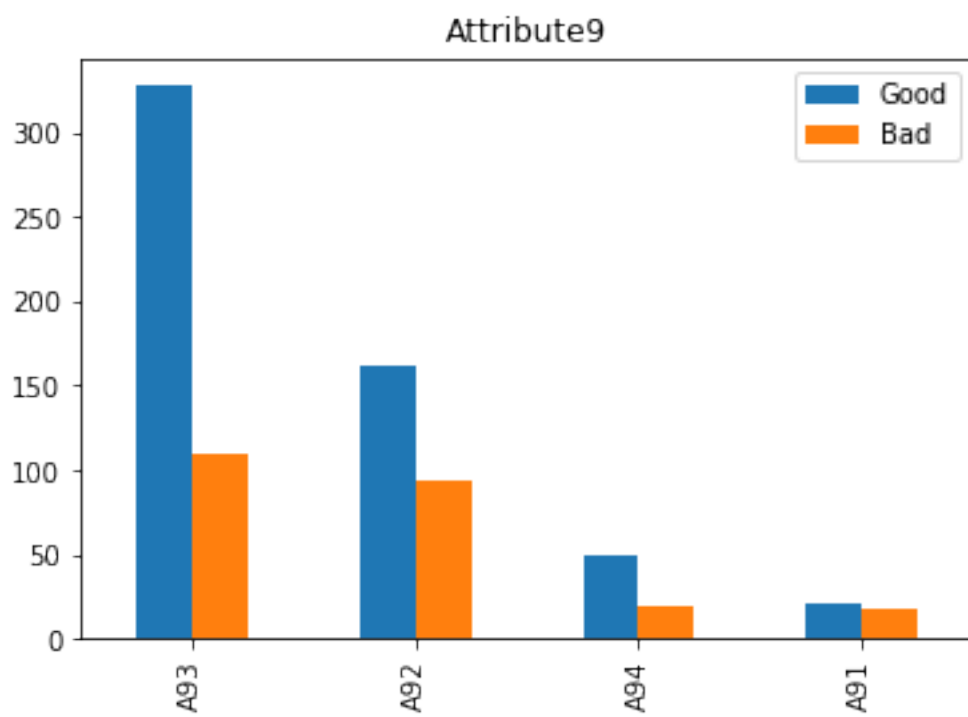

## 2 Categorical Histograms



Observing the histograms, We think that attribute1, attribute3, attribute4 and attribute7 can provide significant information as they have more than 3 values and you can see the difference between the goods and the bads.


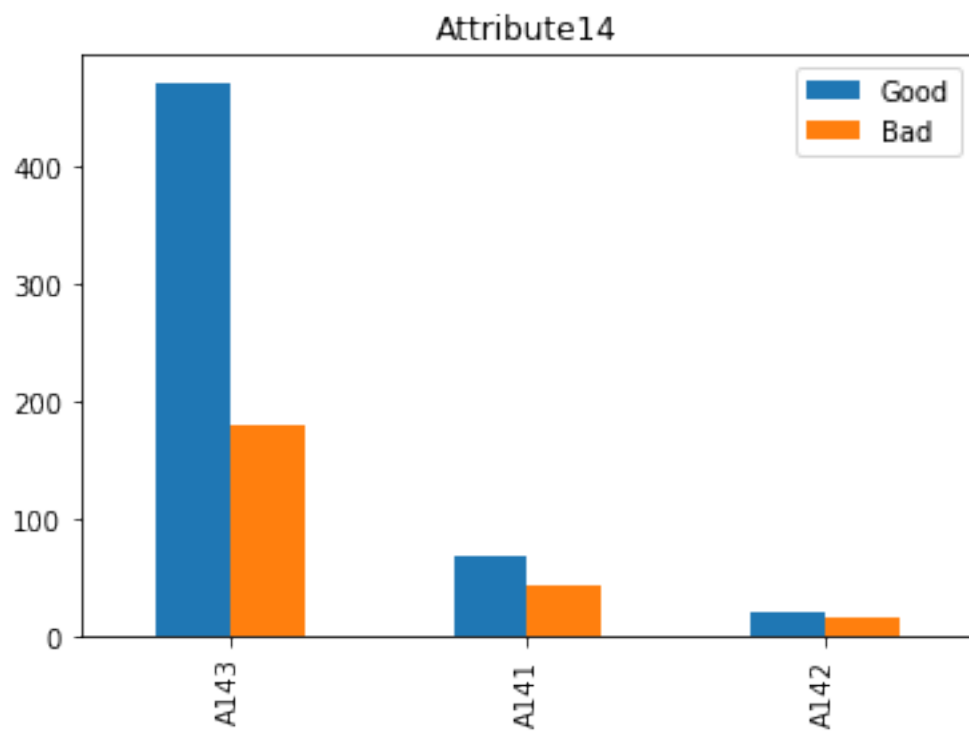
```
In [2]: import matplotlib.pyplot as plt
        cat_cat=["Attribute1", "Attribute3", "Attribute4", "Attribute6", "Attribute7",
                 "Attribute12", "Attribute14", "Attribute15", "Attribute17", "Attribute18"]
        for a_cat in cat_cat:
            new_df=pd.DataFrame({"Good":df_good[a_cat].value_counts(), "Bad":df_bad[a_cat].value_counts()})
            new_df.columns=["Good", "Bad"]
            new_df.plot(kind="bar", title=a_cat)
            plt.show()
```


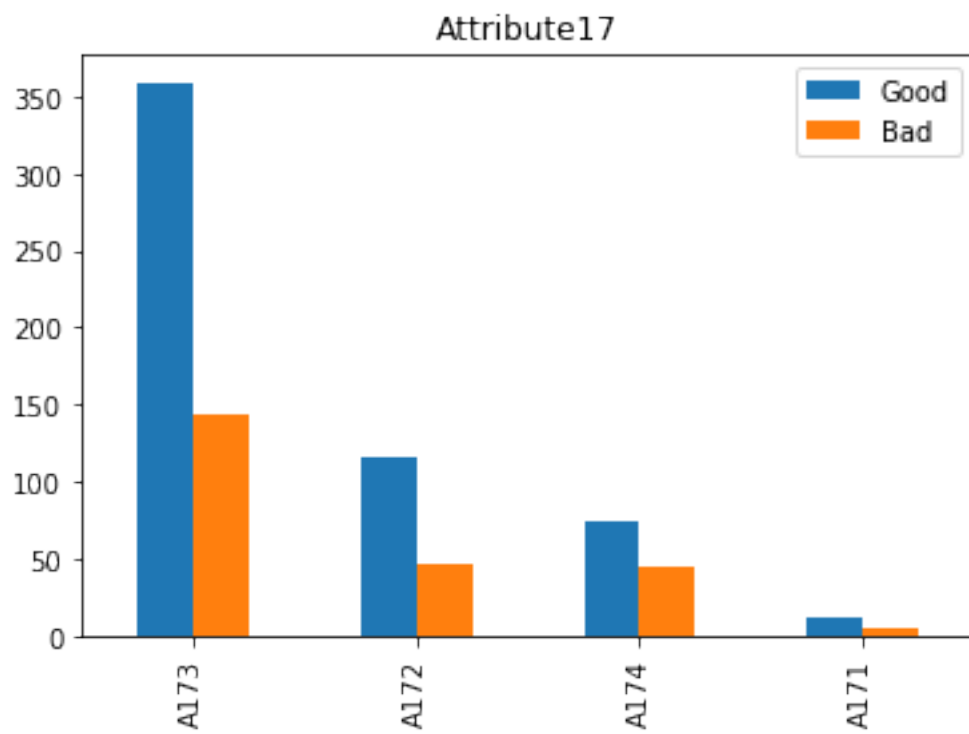



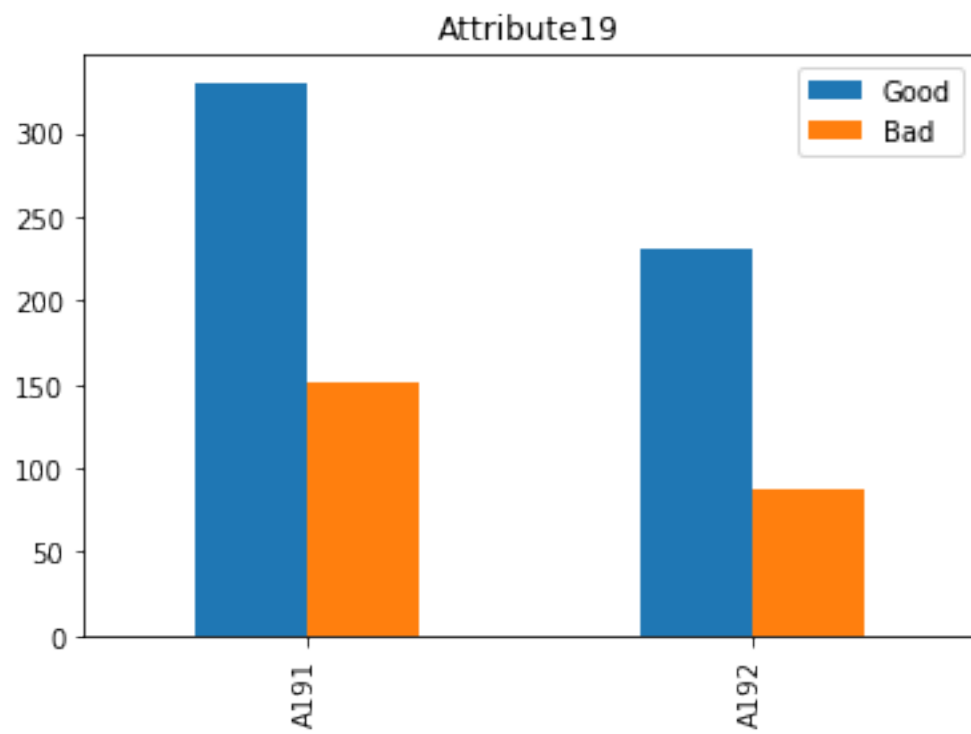










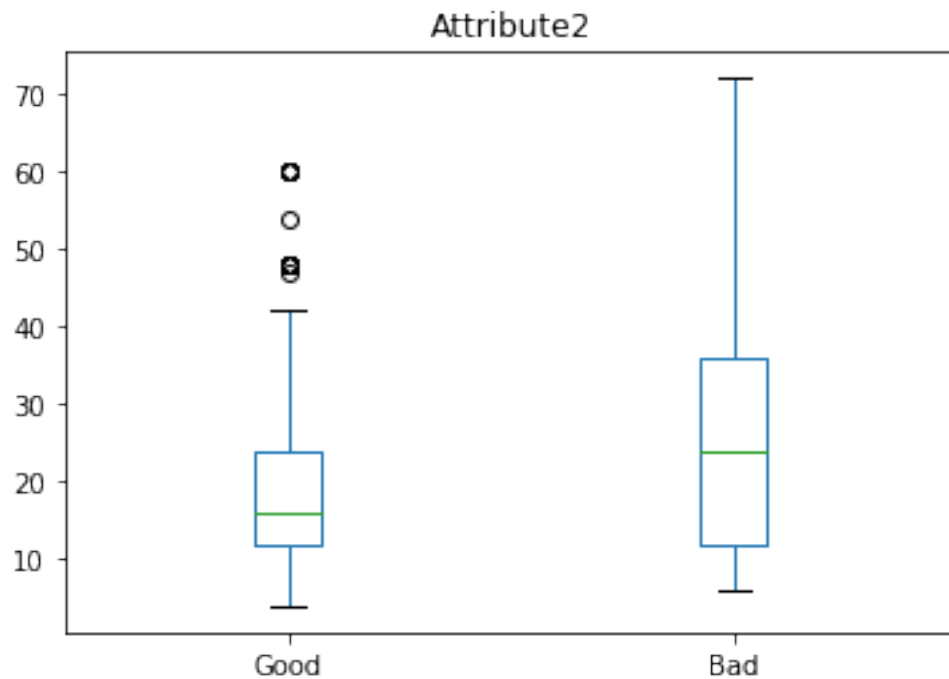



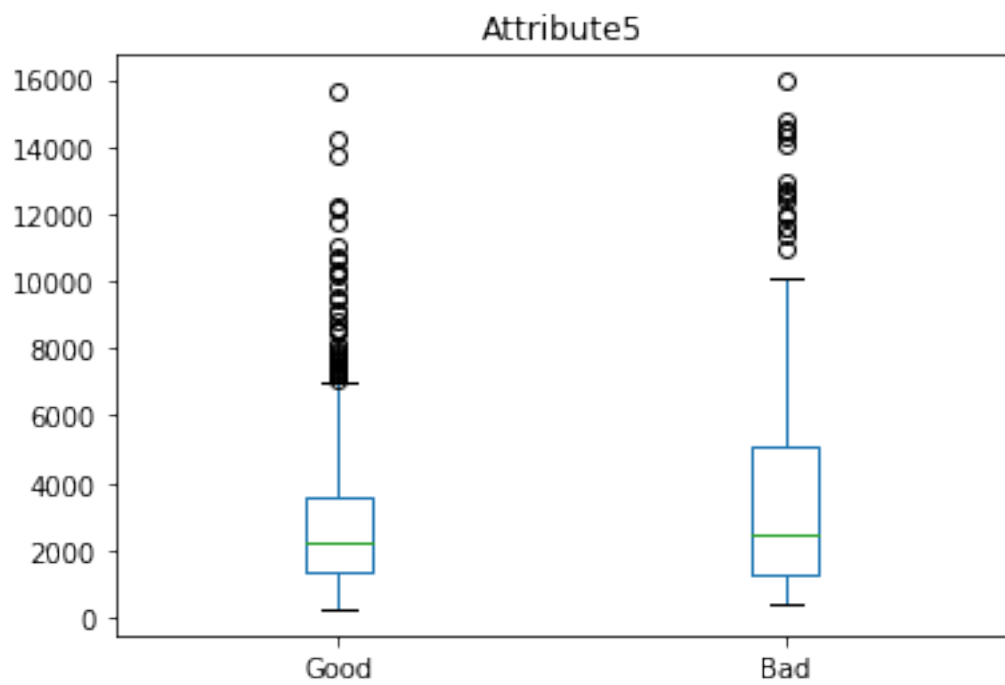



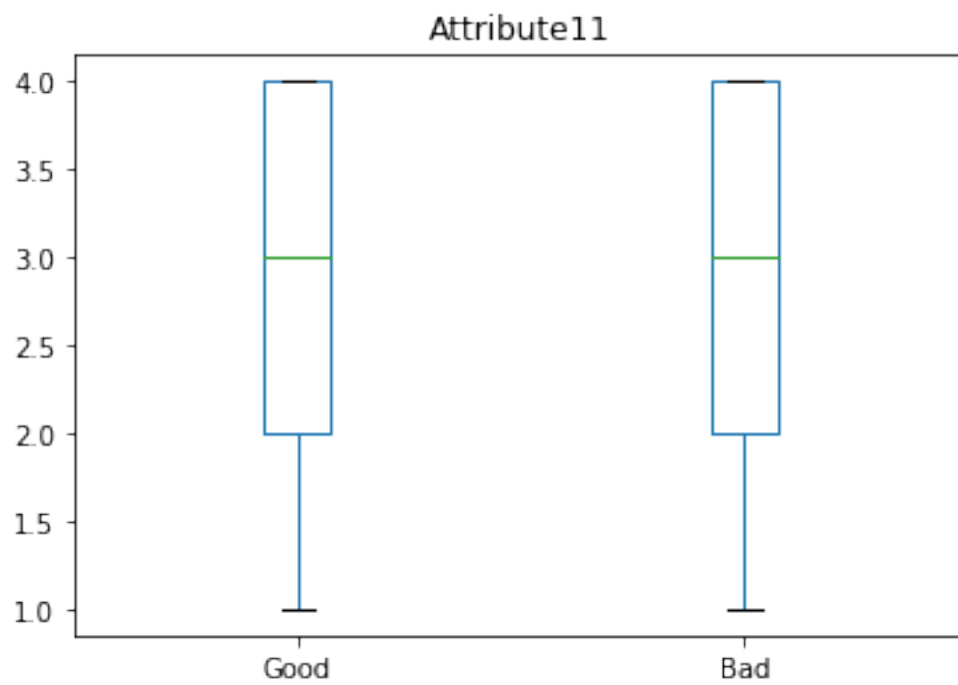
### 3 Numerical box plots


In these box plots, It's easy to say that the attribute2 and attribute5 provide significant information while the other attributes seem to be more unimportant.(Maybe Attribute8 is important?)

```
In [3]: numerical_cat=["Attribute2","Attribute5","Attribute8","Attribute11","Attrib
for a_cat in numerical_cat:
    new_df=pd.DataFrame({"Good":df_good[a_cat],"Bad":df_bad[a_cat]},
                        columns=["Good","Bad"])
    new_df.plot(kind="box",title=a_cat)
    plt.show()
```









## 4 Classifiers

Convert categorical in numeric and cut the id and label columns

```
In [4]: #Copy the dataset
cat_trans=dict()
new_df=df.copy()
for a_cat in cat_cat:
    the_cat_new=pd.Categorical(df[a_cat])
    cat_trans[a_cat]=the_cat_new.categories
    new_df[a_cat]=the_cat_new.codes

#cut the unnecessary
cat_not_needed=["Label","Id"]
cat_needed=df.columns.difference(cat_not_needed)
all_results=dict()
```

## 5 Random Forest

```
In [5]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score
```

```

n_splits=10
kf = KFold(n_splits=n_splits,shuffle=True)
avg_acc=0
for train,test in kf.split(new_df[cat_needed]):
    #getting the correct slices
    my_data=new_df[cat_needed].loc[train]
    labels=new_df["Label"][train]
    #initializeing and fiting the Classifier
    rf_clf=RandomForestClassifier()
    rf_clf.fit(my_data,labels)
    #getting the predictions
    my_pred=rf_clf.predict(new_df[cat_needed].loc[test])
    #the metrics
    acc=accuracy_score(new_df["Label"][test],my_pred)
    avg_acc+=acc/n_splits
print(avg_acc)
all_results["Random Forest"]=avg_acc

```

0.74125

## 6 SVM

```

In [6]: from sklearn import svm
n_splits=10
kf = KFold(n_splits=n_splits,shuffle=True)
avg_acc=0
for train,test in kf.split(new_df[cat_needed]):
    #initializeing and fiting the Classifier
    svm_clas=svm.SVC()
    #getting the correct slices
    my_data=new_df[cat_needed].loc[train]
    labels=new_df["Label"][train]
    svm_clas.fit(my_data,labels)
    #getting the predictions
    my_pred=svm_clas.predict(new_df[cat_needed].loc[test])
    #the metrics
    acc=accuracy_score(new_df["Label"][test],my_pred)
    avg_acc+=acc/n_splits
print(avg_acc)
all_results["SVM"]=avg_acc

```

0.70125

## 7 Naive Bayes (Gaussian NB)

```
In [7]: from sklearn.naive_bayes import GaussianNB
n_splits=10
kf = KFold(n_splits=n_splits,shuffle=True)
avg_acc=0
for train,test in kf.split(new_df[cat_needed]):
    #initializeing and fiting the Classifier
    nbg=GaussianNB()
    #geting the correct slices
    my_data=new_df[cat_needed].loc[train]
    labels=new_df["Label"][train]
    nbg.fit(my_data,labels)
    #geting the predictions
    my_pred=nbg.predict(new_df[cat_needed].loc[test])
    #the metrics
    acc=accuracy_score(new_df["Label"][test],my_pred)
    avg_acc+=acc/n_splits
print(avg_acc)
all_results["Naive Bayes"]=avg_acc
```

0.725

## 8 Write the file of the classifiers

```
In [8]: my_file=open("Evalution_Metric_10fold.csv","w")
for a_clasifier in ["Static Measure","Naive Bayes","Random Forest","SVM"]:
    my_file.write(a_clasifier+"\t")
my_file.write("\n")
my_file.write("Accuracy\t")
for a_clasifier in ["Naive Bayes","Random Forest","SVM"]:
    my_file.write("%.2f\t" % all_results[a_clasifier])
my_file.write("\n")
my_file.write("\n")
my_file.close()
print("Done writting Evalution_Metric_10fold.csv ")
```

Done writting Evalution\_Metric\_10fold.csv

## 9 Calculation of Every Information Gain

```
In [9]: from collections import Counter
from math import log

def my_log(number):
```

```

        if (number<=0):
            return 0.0
        else:
            return log(number)

count_goods=len(df[df["Label"]==1])
count_bads=len(df[df["Label"]==2])
amount=count_goods+count_bads

pos_goods=count_goods/amount
pos_bad=count_bads/amount

df_entropy=-pos_goods*my_log(pos_goods)-pos_bad*my_log(pos_bad)
# attributes that must group
change=["Attribute2","Attribute5","Attribute13"]
# entropy of every feature
entropy_dict=dict()
# information gain of every feature with this dataset
information_gain=dict()
# possibility of every value of every feature
total_pos=dict()
for feature in df[cat_needed].columns:
    counter_good=Counter()
    counter_bad=Counter()
    counter_att=Counter()

    #for one feature
    feature_df=df[feature]
    #if needed, group it
    if(feature in change):
        feature_df=pd.qcut(df[feature],5)
    #count for every value the number of good and bad
    for i in range(amount):
        counter_att[feature_df[i]]+=1
        if df["Label"][i]==1:
            counter_good[feature_df[i]]+=1
        else:
            counter_bad[feature_df[i]]+=1
# possibility of every value from an attribute
    bad=dict()
    good=dict()
    pos=dict()
    for att in counter_att:
        pos[att]=(counter_att[att]/amount)
        good[att]=(counter_good[att]/counter_att[att])
        bad[att]=(counter_bad[att]/counter_att[att])
    # entropy of every value from an attribute
    temp_entropy_dict=dict()

```

```

for att in pos:
    entropy=-good[att]*my_log(good[att])-bad[att]*my_log(bad[att])
    temp_entropy_dict[att]=entropy
total_pos[feature]=pos
#global dict of entropy from every value every feature
entropy_dict[feature]=temp_entropy_dict
#information gain of every feature
for feature in entropy_dict:
    temp_sum=0
    for att in total_pos[feature]:
        temp_sum+=total_pos[feature][att]*entropy_dict[feature][att]
    information_gain[feature]=df_entropy-temp_sum

```

## 10 Random Forest - All accuracies

One feature is taken out in each loop

In [20]: `from collections import OrderedDict`

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score

```

```

def get_accur(cur_df, cat_use):
    n_splits=10
    kf = KFold(n_splits=n_splits, shuffle=True)
    avg_acc=0
    for train, test in kf.split(cur_df[cat_use]):
        #getting the correct slices
        my_data=cur_df[cat_use].loc[train]
        labels=cur_df["Label"][train]
        #initializeing and fitting the Classifier
        rf_clf=RandomForestClassifier(random_state=123)
        rf_clf.fit(my_data, labels)
        #getting the predictions
        my_pred=rf_clf.predict(cur_df[cat_use].loc[test])
        #the metrics
        acc=accuracy_score(cur_df["Label"][test], my_pred)
        avg_acc+=acc/n_splits
    return avg_acc

```

```

information_gain_sorted=[(k, information_gain[k]) for k in sorted(information_gain.keys(), key=lambda k: information_gain[k], reverse=True)]

```

```

new_cat_not_needed=cat_not_needed.copy()
values=[]

```



```

x_data=[]
for num_feature, feature in zip(range(len(information_gain_sorted)-1,0,-1),
    new_cat_not_needed.append(feature[0])
    new_cat_needed=df.columns.difference(new_cat_not_needed)
    cur_accur=get_accur(new_df,new_cat_needed)
    values.insert(0,cur_accur)
    sort_name=feature[0].replace("Attribute","A")
    x_data.insert(0,sort_name+"#"+str(num_feature))

```

## 11 Table with the features and information gains

```

In [21]: # print(information_gain_sorted)
print("-----")
print("| Feature                      | Information Gain      |")
print("-----")
for feat in information_gain_sorted:
    print("| %15s | %15f |" %(feat[0],feat[1]))
print("-----")


```

Feature	Information Gain
Attribute18	0.000090
Attribute11	0.000153
Attribute19	0.000834
Attribute16	0.001661
Attribute17	0.002038
Attribute10	0.003933
Attribute14	0.004881
Attribute8	0.005081
Attribute20	0.005340
Attribute15	0.008054
Attribute13	0.008141
Attribute9	0.008835
Attribute7	0.010084
Attribute12	0.010332
Attribute5	0.010601
Attribute6	0.015387
Attribute4	0.018644
Attribute2	0.022030
Attribute3	0.026263
Attribute1	0.065037

## 12 Plot with the different amount of features

We chose the state that we have 8 attributes (at the point that we take out the attribute 9), because after some executions this point has a good accuracy almost every time.

```
In [22]: import numpy as np
x_np=np.array(range(len(information_gain_sorted)-1,0,-1))
plt.figure(figsize=(20, 8))
plt.xticks(x_np,x_data)
plt.plot(x_np,values,"ro-")
plt.xlabel("Features")
plt.ylabel("Accuracy")
plt.show()
```



## 13 Random Forest For The Test Set

We use some of the features

```
In [23]: #Keep the best categories from the research we did above
best_cat_not_needed=cat_not_needed.copy()
best_cat_not_needed+=["Attribute18","Attribute11","Attribute19","Attribute
best_cat_needed=df.columns.difference(best_cat_not_needed)

In [24]: test_df=pd.read_csv("datasets/test.tsv",delimiter="\t")
# print(test_df)
final_test_df=test_df[best_cat_needed].copy()
for a_cat in cat_cat:
    if(a_cat in best_cat_needed):
        the_cat_new=pd.Categorical(test_df[a_cat],categories=cat_trans[a_c
        final_test_df[a_cat]=the_cat_new.codes

In [25]: best_clf_ever=RandomForestClassifier()
best_clf_ever.fit(new_df[best_cat_needed],new_df["Label"])
my_predictions=best_clf_ever.predict(final_test_df)
```

## 14 Write The Prediction At The File

```
In [26]: my_file=open("testSet_Predictions.csv", "w")
        for a_clasifier in ["Client_ID", "Predicted_Label"]:
            my_file.write(a_clasifier+"\t")
        my_file.write("\n")
        for i in range(len(my_predictions)):
            my_file.write(str(test_df["Id"][i])+"\t")
            if(my_predictions[i]==1):
                my_file.write("Good\n")
            else:
                my_file.write("Bad\n")

        my_file.write("\n")
        my_file.close()
        print("Done writting testSet_Predictions.csv")
```

Done writting testSet\_Predictions.csv

```
In [ ]:
```